

Assignment 4

SMS SPAM Classification

Assignment Date	08 November 2022
Student Name	S.Karthikeyan
Student Register Number	620619106012
Maximum Marks	2

1. Downloading the Dataset

<https://www.kaggle.com/code/kredy10/simple-lstm-for-text-classification/data>

2. Import required library

```
import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder
from keras.models import Model
from keras.layers import LSTM, Activation, Dense, Dropout, Input, Embedding
from keras.optimizers import RMSprop
from keras.preprocessing.text import Tokenizer
from keras_preprocessing import sequence
from keras.utils import to_categorical
from keras.models import load_model
```

```
[1] import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder
from keras.models import Model
from keras.layers import LSTM, Activation, Dense, Dropout, Input, Embedding
from keras.optimizers import RMSprop
from keras.preprocessing.text import Tokenizer
from keras_preprocessing import sequence
from keras.utils import to_categorical
from keras.models import load_model
```

3. Read Dataset and do preprocessing

```
df = pd.read_csv('/content/drive/MyDrive/Assignment 3/spam.csv',  
delimiter=',',encoding='latin-1')  
df.head()
```

```
[3] df = pd.read_csv('/content/drive/MyDrive/Assignment 3/spam.csv',delimiter=',',encoding='latin-1')
```

```
[4] df.head()
```

	v1	v2	Unnamed: 2	Unnamed: 3	Unnamed: 4
0	ham	Go until jurong point, crazy.. Available only ...	NaN	NaN	NaN
1	ham	Ok lar... Joking wif u oni...	NaN	NaN	NaN
2	spam	Free entry in 2 a wkly comp to win FA Cup fina...	NaN	NaN	NaN
3	ham	U dun say so early hor... U c already then say...	NaN	NaN	NaN
4	ham	Nah I don't think he goes to usf, he lives aro...	NaN	NaN	NaN

```
df.drop(['Unnamed: 2', 'Unnamed: 3', 'Unnamed: 4'],axis=1,inplace=True) #d  
ropping unwanted columns  
df.info()
```

```
[5] df.drop(['Unnamed: 2', 'Unnamed: 3', 'Unnamed: 4'],axis=1,inplace=True) #dropping unwanted columns
```

```
[6] df.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 5572 entries, 0 to 5571  
Data columns (total 2 columns):  
#   Column  Non-Null Count  Dtype  
---  ---  
0    v1      5572 non-null       object  
1    v2      5572 non-null       object  
dtypes: object(2)  
memory usage: 87.2+ KB
```

```
# Count of Spam and Ham values
```

```
df.groupby(['v1']).size()
```

```
[7] # Count of Spam and Ham values  
df.groupby(['v1']).size()
```

```
v1  
ham    4825  
spam    747  
dtype: int64
```

```
# Label Encoding target column
```

```
X = df.v2
```

```
Y = df.v1
```

```
le = LabelEncoder()
```

```
Y = le.fit_transform(Y)
```

```
Y = Y.reshape(-1,1)
```

```
# Test and train split
```

```
X_train,X_test,Y_train,Y_test = train_test_split(X,Y,test_size=0  
.15)
```

```
# Tokenisation function
```

```
max_words = 1000
```

```

max_len = 150
tok = Tokenizer(num_words=max_words)
tok.fit_on_texts(X_train)
sequences = tok.texts_to_sequences(X_train)

sequences_matrix = sequence.pad_sequences(sequences,maxlen=max_len)

```

```

[8] # Label Encoding target column
X = df.v2
Y = df.v1
le = LabelEncoder()
Y = le.fit_transform(Y)
Y = Y.reshape(-1,1)

```

```

[9] # Test and train split
X_train,X_test,Y_train,Y_test = train_test_split(X,Y,test_size=0.15)

```

```

[10] # Tokenisation function
max_words = 1000
max_len = 150
tok = Tokenizer(num_words=max_words)
tok.fit_on_texts(X_train)
sequences = tok.texts_to_sequences(X_train)

sequences_matrix = sequence.pad_sequences(sequences,maxlen=max_len)

```

4. Create Model

```

# Creating LSTM model
inputs = Input(name='InputLayer', shape=[max_len])

```

```

[11] # Creating LSTM model
inputs = Input(name='InputLayer', shape=[max_len])

```

5. Add Layers (LSTM, Dense – (Hidden Layers), Output)

```

layer = Embedding(max_words, 50, input_length=max_len)(inputs)
layer = LSTM(64)(layer)
layer = Dense(256, name='FullyConnectedLayer1')(layer)
layer = Activation('relu')(layer)
layer = Dropout(0.5)(layer)
layer = Dense(1, name='OutputLayer')(layer)
layer = Activation('sigmoid')(layer)

```

```
[12] layer = Embedding(max_words,50,input_length=max_len)(inputs)
      layer = LSTM(64)(layer)
      layer = Dense(256,name='FullyConnectedLayer1')(layer)
      layer = Activation('relu')(layer)
      layer = Dropout(0.5)(layer)
      layer = Dense(1,name='OutputLayer')(layer)
      layer = Activation('sigmoid')(layer)
```

6. Compile the model

```
model = Model(inputs=inputs,outputs=layer)
model.summary()
```

```
[13] model = Model(inputs=inputs,outputs=layer)
```

```
[14] model.summary()
```

Model: "model"

Layer (type)	Output Shape	Param #
=====		
InputLayer (InputLayer)	[(None, 150)]	0
embedding (Embedding)	(None, 150, 50)	50000
lstm (LSTM)	(None, 64)	29440
FullyConnectedLayer1 (Dense)	(None, 256)	16640
activation (Activation)	(None, 256)	0
dropout (Dropout)	(None, 256)	0
OutputLayer (Dense)	(None, 1)	257
activation_1 (Activation)	(None, 1)	0

```
=====
Total params: 96,337
Trainable params: 96,337
Non-trainable params: 0
```

```
model.compile(loss='binary_crossentropy',optimizer=RMSprop(),metrics=['accuracy'])
```

```
[15] model.compile(loss='binary_crossentropy',optimizer=RMSprop(),metrics=['accuracy'])
```

7. Fit the Model

```
model.fit(sequences_matrix,Y_train,batch_size=128,epochs=10,
          validation_split=0.2)
```

```
[16] model.fit(sequences_matrix,Y_train,batch_size=128,epochs=10,
              validation_split=0.2)

Epoch 1/10
30/30 [=====] - 13s 288ms/step - loss: 0.3222 - accuracy: 0.8801 - val_loss: 0.1518 - val_accuracy: 0.9515
Epoch 2/10
30/30 [=====] - 8s 267ms/step - loss: 0.0850 - accuracy: 0.9791 - val_loss: 0.0584 - val_accuracy: 0.9863
Epoch 3/10
30/30 [=====] - 11s 363ms/step - loss: 0.0390 - accuracy: 0.9900 - val_loss: 0.0530 - val_accuracy: 0.9884
Epoch 4/10
30/30 [=====] - 8s 266ms/step - loss: 0.0293 - accuracy: 0.9910 - val_loss: 0.0516 - val_accuracy: 0.9895
Epoch 5/10
30/30 [=====] - 8s 268ms/step - loss: 0.0230 - accuracy: 0.9926 - val_loss: 0.0677 - val_accuracy: 0.9852
Epoch 6/10
30/30 [=====] - 8s 268ms/step - loss: 0.0164 - accuracy: 0.9947 - val_loss: 0.0723 - val_accuracy: 0.9852
Epoch 7/10
30/30 [=====] - 8s 268ms/step - loss: 0.0122 - accuracy: 0.9963 - val_loss: 0.0754 - val_accuracy: 0.9852
Epoch 8/10
30/30 [=====] - 8s 265ms/step - loss: 0.0109 - accuracy: 0.9958 - val_loss: 0.0904 - val_accuracy: 0.9852
Epoch 9/10
30/30 [=====] - 8s 265ms/step - loss: 0.0060 - accuracy: 0.9979 - val_loss: 0.1023 - val_accuracy: 0.9852
Epoch 10/10
30/30 [=====] - 8s 267ms/step - loss: 0.0042 - accuracy: 0.9987 - val_loss: 0.1120 - val_accuracy: 0.9873
<keras.callbacks.History at 0x7fa615cafad0>
```

8. Save the Model

```
model.save("model_1")
```

```
[17] model.save("model_1")

WARNING:absl:Function `wrapped_model` contains input name(s) InputLayer with unsupported characters which will be renamed to inputlayer in the SavedModel.
WARNING:absl:Found untraced functions such as lstm_cell_layer_call_fn, lstm_cell_layer_call_and_return_conditional_losses while saving (showing 2 of 2). These fu
```

9. Test the Model

```
test_sequences = tok.texts_to_sequences(X_test)
```

```
test_sequences_matrix = sequence.pad_sequences(test_sequences,maxlen=max_len)
```

```
accuracy = model.evaluate(test_sequences_matrix,Y_test)
```

```
[18] test_sequences = tok.texts_to_sequences(X_test)
```

```
[19] test_sequences_matrix = sequence.pad_sequences(test_sequences,maxlen=max_len)
```

```
[20] accuracy = model.evaluate(test_sequences_matrix,Y_test)
```

```
27/27 [=====] - 1s 34ms/step - loss: 0.0923 - accuracy: 0.9809
```

```
print('Accuracy: {:.3f}'.format(accuracy[1]))
```

```
[21] print('Accuracy: {:.3f}'.format(accuracy[1]))
```

Accuracy: 0.981

```
y_pred = model.predict(test_sequences_matrix)
```

```
[22] y_pred = model.predict(test_sequences_matrix)
```

```
27/27 [=====] - 2s 34ms/step
```

```
print(y_pred[25:40].round(3))
```

```
[23] print(y_pred[25:40].round(3))
```

```
print(Y_test[25:40])
```

```
[24] print(Y_test[25:40])
```